

PROMPT ENGINEERING ASSIGNMENT

The Prompt Ladder.

Six runs, one layer at a time. From “Help me with marketing” to a brief that writes itself — with the version that made things worse left in.

TRACK Marketing · SaaS launch brief

METHOD One layer per version, side-by-side compare

DATE September 2026

Introduction & Methodology

This document tracks the evolution of a single weak prompt (“*Help me with marketing.*”) across five versions. Each version adds exactly one layer — a clearer goal, real context, a defined audience, a specified output format, or constraints — chosen to attack the previous output’s biggest weakness. Every version was run, and its output is compared side-by-side with the previous one. Four short notes follow each version: what changed in the prompt, what actually improved in the output, what still failed, and what I would try next.

Method. One change per version, no bundling. The four notes describe *changes in the output*, not just the prompt. At least one version is an honest “this made it worse” moment — if every version improved smoothly, I wasn’t looking hard enough.

Baseline — No layer (the embarrassing starting point)

Prompt:

Help me with marketing.

Model output (excerpt):

“Sure! Marketing is a broad field that involves promoting products or services to target audiences. To help you effectively, I’d need to know more about what you’re marketing. Are you looking for help with: digital marketing (social media, email, SEO), content marketing, brand strategy, paid advertising, or market research? Once you let me know your specific needs, I can provide tailored advice!”

What changed	Nothing. This is the baseline — the prompt I would have typed a month ago.
What improved	N/A — the output is a clarifying question, not a deliverable. The model is asking me to do its job.
What still failed	No brief, no recommendation, no structure. Five bulleted sub-fields of marketing, none of them actionable.
What to try next	Add a clearer goal so the model writes a brief instead of asking me which kind of marketing I want.

Versions 1 & 2

Version 1 — Layer: clearer goal

Prompt:

Help me write a marketing brief for launching our new SaaS product.

Model output (excerpt):

"Happy to help! A marketing brief for a SaaS launch typically includes: (1) Target Audience — who you're selling to; (2) Value Proposition — what makes your product unique; (3) Channels — where you'll reach your audience; (4) Budget & Timeline — when and how much; (5) KPIs — how you'll measure success. Could you share details about your product, target customer, and budget? Once I have that, I can draft a tailored brief."

What changed	Specified the deliverable ("write a marketing brief") and the context ("launching our new SaaS product").
What improved	Output is now structured — lists the 5 sections a brief should contain, instead of asking which sub-field of marketing I care about.
What still failed	Model still ends by asking me for product details. It's an outline of a brief, not a brief.
What to try next	Feed real context (product, audience, pricing) so the model can actually draft, not just enumerate.

Version 2 — Layer: real context

Prompt:

Help me write a marketing brief for launching our new SaaS product.

Context:

- Product: Pulse, a Slack-native standup bot for engineering teams (50-200 devs)
- Pricing: \$6/user/month, 14-day free trial, no credit card
- Launching: 6 weeks, producthunt.com launch planned
- Channels we can use: Product Hunt, engineering subreddits, dev Twitter, Slack discover
- We have \$0 paid budget, only owned + community channels

Model output (excerpt):

"Marketing Brief — Pulse Launch. Objective: Drive 500 trial signups in launch week via community-led channels. **Audience:** Engineering managers at 50-200 dev shops who already run async standups in Slack. **Key Message:** 'Standups that run themselves — Pulse posts the prompt, collects updates, summarizes for you.' **Channels:** Product Hunt (launch-day teaser, maker comment thread, founder AMA 10am PT); r/ExperiencedDevs (case-study post, 'How we cut standup time by 40%'); dev Twitter (5-day thread on standup anti-patterns); Slack App Discovery (listing copy refresh). **KPIs:** Trial signups, activation rate (first standup run within 24h), Slack listing CTR. **Risks:** Product Hunt algorithm rewards early momentum — need 3-5 committed upvoters at midnight PT."

What changed	Added real context — product, audience, pricing, timing, channel constraints, budget.
What improved	Output is now an actual brief, not an outline. Specific tactics, real numbers, named channels, risks called out.
What still failed	Tone is a bit generic — reads like a competent agency template, not tailored to a specific reader. Risks section is thin.
What to try next	Specify the audience for the brief itself (who will read it) to sharpen the tone.

Versions 3 & 4

Version 3 — Layer: defined audience — HONEST MOMENT (output got worse)

Prompt:

Help me write a marketing brief for launching our new SaaS product.
Context: [same as V2]
Audience for this brief: Senior PMs and marketing leads who already know the basics. Skip explanations of what a brief is or what KPIs are.

Model output (excerpt) — drier than V2:

“Pulse Launch Brief. Objective: 500 trial signups, launch week. Audience: Eng managers, 50-200 dev teams, async standup adopters. Channels: Product Hunt, r/ExperiencedDevs, dev Twitter, Slack Discovery. KPIs: Trials, activation, listing CTR. Risks: PH momentum dependency.”

What changed	Added an explicit audience for the brief itself — “senior PMs and marketing leads who already know the basics.”
What improved	Honestly, not much. Output got shorter and drier, but lost the tactical specificity (the “maker comment thread at 10am PT” line, the “5-day thread series”) that made V2 actually useful.
What still failed	By telling the model “skip the basics,” it over-corrected and skipped the actionable stuff too. The brief is now too thin to hand to anyone.
What to try next	Restore tactical detail by specifying an output format that forces structure — sections with minimum content per section.

Version 4 — Layer: specified output format

Prompt:

Help me write a marketing brief for launching our new SaaS product.
Context: [same as V2]
Audience for this brief: Senior PMs and marketing leads who already know the basics.

Output format – use exactly these sections, in this order:

1. Objective (1-2 sentences, with a number)
2. Audience (1 paragraph: who, what they already use, what they hate)
3. Channels (1 sub-bullet per channel with the specific tactic for THIS launch)
4. Key Message (1 sentence the team can put on a slide)
5. KPIs (each with a 30-day target)
6. Risks & Dependencies (with a mitigation note)

Model output (excerpt):

“1. Objective. 500 trial signups, launch week (Aug 12-18), \$0 paid spend. 2. Audience. Eng managers at 50-200 dev teams running async standups in Slack but hate the manual nagging. Tried Geekbot/Standuply, found them bloated. Read r/ExperiencedDevs, distrust ‘platforms.’ 3. Channels. PH: launch 12:01am PT, founder AMA 10am-12pm. r/ExperiencedDevs: case-study post, no direct CTA. Dev Twitter: 5-day anti-patterns thread, day 5 soft launch. Slack Discovery: A/B test description copy. 4. Key Message. ‘Standups that run themselves.’ 5. KPIs. Trials 500 (wk1), 1,500 (cumulative day 30). Activation 60%. Paid conversion 8%. 6. Risks. PH algorithm penalty if <100 upvotes by 9am → 50 pre-lined upvoters. Reddit self-promo removal → non-team member posts, case-study lead.”

What changed	Specified the exact section structure with a minimum content hint per section.
What improved	Format constraint forced the model to fill in detail per section instead of compressing. Got back the tactical specificity V3 lost, plus new things V2 didn’t have (30-day conversion target, mod-removal mitigation).
What still failed	Some numbers feel plausible but invented (the “8% conversion” line). No way to tell which are real vs guessed.
What to try next	Add constraints that force the model to flag what it doesn’t know instead of inventing.

Version 5 & Comparison

Version 5 — Layer: constraints (mark assumptions)

Prompt:

Help me write a marketing brief for launching our new SaaS product.

Context: [same as V2]

Audience: [same as V3]

Output format: [same as V4]

Constraints:

- Maximum 400 words total.
- Do not invent pricing, conversion rates, or competitive numbers. If you don't know, write [ASSUMPTION] and what kind of input would resolve it.
- Each KPI must be measurable within 30 days.
- Mark any tactic you're guessing works (no internal data) with [GUESS].

Model output (excerpt):

1. Objective. 500 trial signups, launch week (Aug 12-18), \$0 paid. **2. Audience.** Eng managers, 50-200 dev teams, async standups in Slack. Tried Geekbot/Standuply, found them bloated. **3. Channels.** PH: 12:01am PT, AMA 10am-12pm. [GUESS: 5 pre-committed upvoters enough for top-10]. r/ExperiencedDevs: case-study post, no CTA. [GUESS: case-study angle beats launch announcement]. Dev Twitter: 5-day thread. Slack Discovery: A/B test copy. **4. Key Message.** 'Standups that run themselves.' **5. KPIs.** Trials 500/1,500. Activation 60% [ASSUMPTION: from beta data, need confirmation]. Paid conversion [ASSUMPTION: SaaS trial-to-paid is 5-15%, need our own historical]. **6. Risks.** PH <100 upvotes by 9am → 50 pre-lined. Reddit self-promo → non-team member posts."

What changed Added constraints — word cap, no invented numbers, mark guesses and assumptions explicitly.

What improved The output now tells you what it doesn't know. The fake conversion-rate line is gone (replaced with an honest [ASSUMPTION] tag) and tactics flag which are guesses vs known. The brief is now actually trustworthy enough to hand to someone.

What still failed Word count forced some loss of nuance in section 2 (audience paragraph got thinner). The [GUESS] tags could be more specific about what evidence would update them.

What to try next This is good enough to ship. The next iteration would add examples of "good brief" to anchor quality, but the marginal gain is small.

Comparison Matrix

All six runs side-by-side. The V3 row is honestly flagged — the audience layer made the output thinner, not sharper.

Version	Layer Added	What Improved in Output	What Still Failed
Baseline	—	Nothing — model asks a clarifying question	No deliverable produced
V1	Clearer goal	Lists the 5 sections a brief should contain	Still asks for input; outline, not brief
V2	Real context	Drafts an actual brief with tactics, channels, numbers	Tone generic; risks section thin
V3	Defined audience	Got shorter — lost tactical specificity	Over-corrected; too thin to hand to anyone
V4	Output format	Restored detail; added 30-day targets + mitigations	Some numbers feel invented (8% conversion)
V5	Constraints	Flags [ASSUMPTION] / [GUESS] tags; honest about unknowns	Some audience nuance lost to word cap

Table 1. Side-by-side comparison. The honest moment (V3) is highlighted in muted red.

Final Reusable Prompt & Reflection

The Final Prompt

Cleaned up so someone else on the marketing track could take it and use it without me in the room. Replace the bracketed fields with your own product details.

Help me write a marketing brief for launching our new SaaS product.

Context:

- Product: [name + 1-line description]
- Pricing: [price + trial terms]
- Launching: [timeline + flagship launch channel]
- Channels we can use: [list, with constraints like "\$0 paid"]
- Audience for the product: [who, what they already use, what they hate]

Audience for this brief: Senior PMs and marketing leads who already know the basics. Skip explanations of what a brief is or what KPIs are.

Output format – use exactly these sections, in this order:

1. Objective (1-2 sentences, with a number)
2. Audience (1 paragraph: who, what they already use, what they hate)
3. Channels (1 sub-bullet per channel with the specific tactic for THIS launch)
4. Key Message (1 sentence the team can put on a slide)
5. KPIs (each with a 30-day target)
6. Risks & Dependencies (with a mitigation note)

Constraints:

- Maximum 400 words total.
- Do not invent pricing, conversion rates, or competitive numbers. If you don't know, write [ASSUMPTION] and what kind of input would resolve it.
- Each KPI must be measurable within 30 days.
- Mark any tactic you're guessing works (no internal data) with [GUESS].

Reflection

The biggest surprise was Version 3. I assumed “tell it who the reader is” would be the highest-leverage move — every prompt engineering guide says audience matters, and usually it does. But for a brief that was already short, telling the model “skip the basics” gave it permission to skip the parts that mattered too. The tactical specificity — the “maker comment thread at 10am PT,” the “5-day anti-patterns thread” — evaporated. The lesson: an audience layer sharpens a verbose output, but it cannot rescue a thin one. Audience works when there is enough material to filter.

The second surprise was how much Version 5's [ASSUMPTION] and [GUESS] tags changed the *felt* quality of the output. The brief didn't get smarter — it got more honest. And honest is what makes a brief trustworthy enough to hand to someone else. That is the layer I would keep first if I had to drop the others. A brilliant brief with invented numbers is worse than a mediocre brief that admits what it doesn't know, because the former gets you confidently wrong and the latter invites the right follow-up question.

What I would try next, given another round: **examples of what good looks like** — one real brief, ideally from a comparable launch. The format layer told the model the shape; an example would tell it the bar. That is the iteration I expect would compress the gap between V5 and a genuinely senior writer's first draft more than any other single change.